

42: The answer to **life**, the **universe**, and **everything**
offensive security

Who Are We?



Will Pearce
moo_hax

Nick Landers
monoxgas




Work at **Silent Break Security**
Research, Dev, Training, Ops, etc.



What is Machine Learning?

- Lots of magic that:
gets investors, goes public, makes \$
- Rebranded if/else statements
- Data Sheet Keywords:
 - “Analyzes millions of X and adapts”
 - “High degree of confidence”
 - “Next Generation”
 - “Scientists”
 - “Central Cortex”
 - “Neural Network”



What **really** is ML?

- Set of techniques that aim to:
model a problem mathematically
- Stats + Maths + Computers
- Very old discipline – 1950's (IBM)
- Predictions without explicit programming
- **Growing fast:** computing power, data aggregation, etc.

... Still magic ... But **mostly math**



Why Do We Care?

- It's coming to a field near you
 - No longer a math problem, it's an engineering problem
 - Our future will be fought with ML
- It can be really **really** awesome
 - Building relationships in non-congruent data
 - Bring out operator 6th senses
 - Crush huge amounts of data faster than humans
 - Can be as complex or as simple as you want to make it
 - Optimizing a traditionally “manual” line of work



What is “Offensive” ML?

“The application of ML to offensive security problems”

- Abusing control relationships * [**Neo4j/Kegra**]
- Obfuscating C2 as English [**Markov Obfuscate**]
- Detecting sandbox environments * [**Deep Drop**]
- Improving social media phishing [**SNAP_R**]
- Faster password guessing [**PassGan**]
- Metasploit exploit selection [**Deep Exploit**]
- Stealing models for evasion * [**Adversarial ML**]
- Automating timing attacks [**ParzelSec**]

* - to be discussed



Collision Course



Starting with ML



“ML [literally anything] tutorial”

“Detecting Cats in Images with OpenCV”

“Auto-Generating Clickbait with RNNs”

<https://github.com/ujjwalkarn/Machine-Learning-Tutorials/> (320+ links)

<https://sgfin.github.io/learning-resources/> (200+ links)

<https://github.com/josephmisiti/awesome-machine-learning> (200+ links)

<https://github.com/awesomedata/awesome-public-datasets> (410+ links)

Starting with ML

1. Data is everywhere, what ~~can~~ should we collect?
2. What data is used by a human to solve/perform X?
3. Process the data and extract useful features.
4. Download Python + [ML stuff]. (Might need some GPUs)
 - [NumPy / Pandas](#) – Data processing and matrices
 - [SciKit-Learn](#) – Data analytics + Basic ML
 - [TensorFlow](#) – Full blown ML framework from Google
 - [Keras](#) – High-level wrapper for TensorFlow (also CNTK, Theano)
5. Write a 10 line script and ML your heart out.



Sandbox Detection – Case Study

- Sandboxes are a dangerous place
 - Popularity is rising
 - Preventing analysis is a priority
- Lots of current detection strategies
 - Enumerating host information
 - Automated behavior indicators
 - Network anomalies – server and client
 - ...

Need detection using minimal information





ML Concepts

Data – The raw information we gather

PID	User	Process
1	NT AUTHORITY\SYSTEM	smss.exe
236	NT AUTHORITY\SYSTEM	csrss.exe
120	NT AUTHORITY\SYSTEM	winlogon.exe
492	Admin-PC\Admin	explorer.exe
940	Admin-PC\Admin	procmon.exe
680	Admin-PC\Admin	dllhost.exe
772	Admin-PC\Admin	winword.exe





ML Concepts

Features – How we represent data to an algorithm

PID	User	Process		
1	NT AUTHORITY\SYSTEM	smss.exe		
236	NT AUTHORITY\SYSTEM	csrss.exe		
120	NT AUTHORITY\SYSTEM	winlogon.exe		
492	Admin-PC\Admin	explorer.exe		
940	Admin-PC\Admin	procmon.exe		
680	Admin-PC\Admin	dllhost.exe		
772	Admin-PC\Admin	winword.exe		
			Feature	Value
			bad_user:	1
			sysinternals:	1
			domain_member:	false



ML Concepts

Features – How we represent data to an algorithm

PID	User	Process		
1	NT AUTHORITY\SYSTEM	smss.exe		
36	NT AUTHORITY\SYSTEM	csrss.exe		
312	NT AUTHORITY\SYSTEM	winlogon.exe		
444	ACME\arthur.dent	explorer.exe		
452	ACME\arthur.dent	winword.exe		
1972	ACME\arthur.dent	chrome.exe		
2928	ACME\arthur.dent	chrome.exe		
			Feature	Value
			proc_count:	7
			average_pid	906
			compression:	85%
			proc_per_user:	3.5






ML Concepts

Inputs – Our features for each sample

```
data = np.array([  
    [33, 4, 8.25],  
    [157, 1, 157],  
    [195, 1, 195],  
    [30, 4, 7.5],  
    [34, 4, 8.5],  
    [84, 1, 84]  
])
```





ML Concepts

Label – The thing we are predicting *

```
data = np.array([
    [33, 4, 8.25, 1],
    [157, 1, 157, 0],
    [195, 1, 195, 0],
    [30, 4, 7.5, 1],
    [34, 4, 8.5, 1],
    [84, 1, 84, 0]
])
```

0 = safe
1 = sandbox

* - we manually label data for training





ML Concepts

Classification – The strategy for assessing outputs

Binary: Grouping into 2 fixed categories

Multi-Label: Grouping into N categories

Regression: Targeting a continuous variable value
(1 – 999)

* - we usually have labeled data for training



ML Concepts


Node – Smallest unit to carry an activation



* Calculated using an
activation function
(like sigmoid)


$$\frac{1}{1 + e^{-x}}$$





ML Concepts

Layer(s) – A parallel group of nodes



input



hidden




output





ML Concepts

Network – A few layers strapped together *



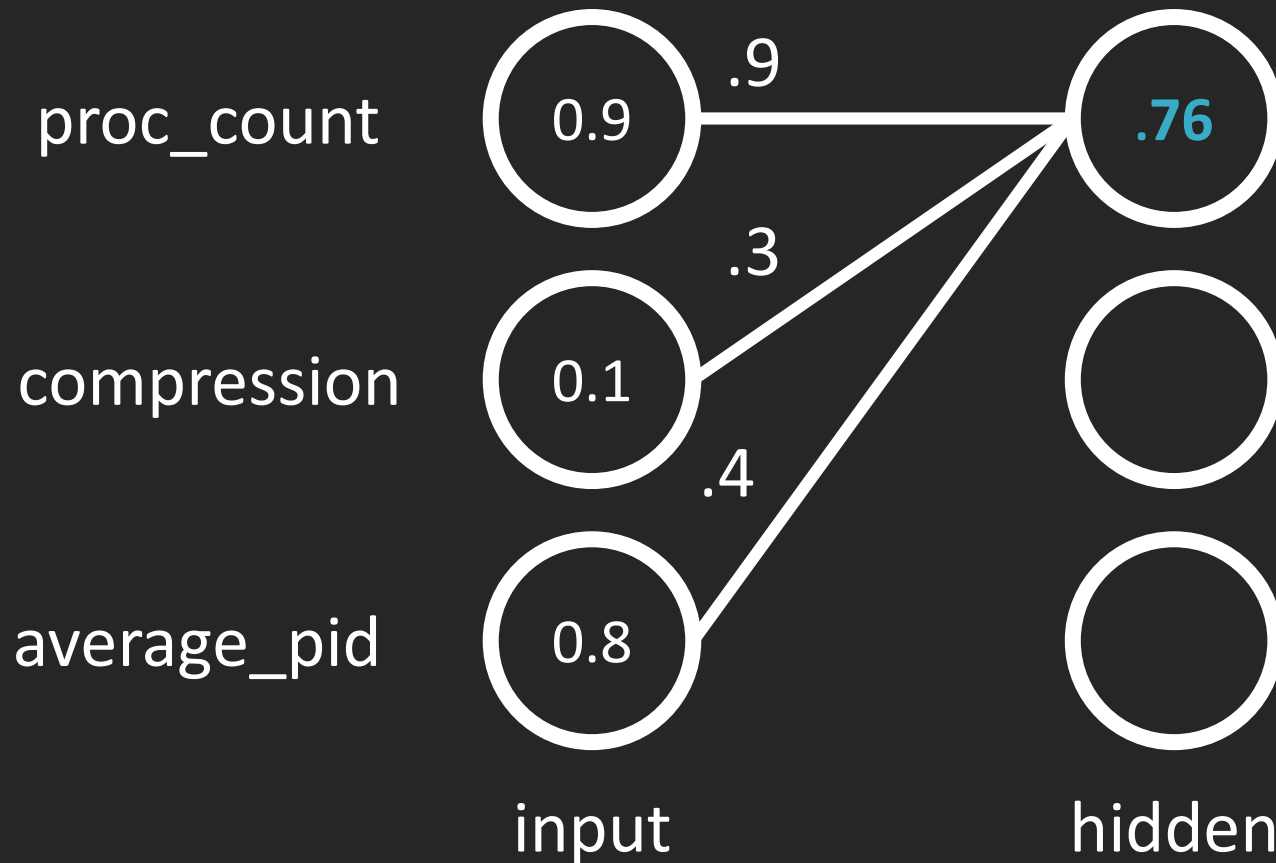
* - A “Neural Network” isn’t the only type of network





ML Concepts


Feed Forward – calculate the error of the network





ML Concepts

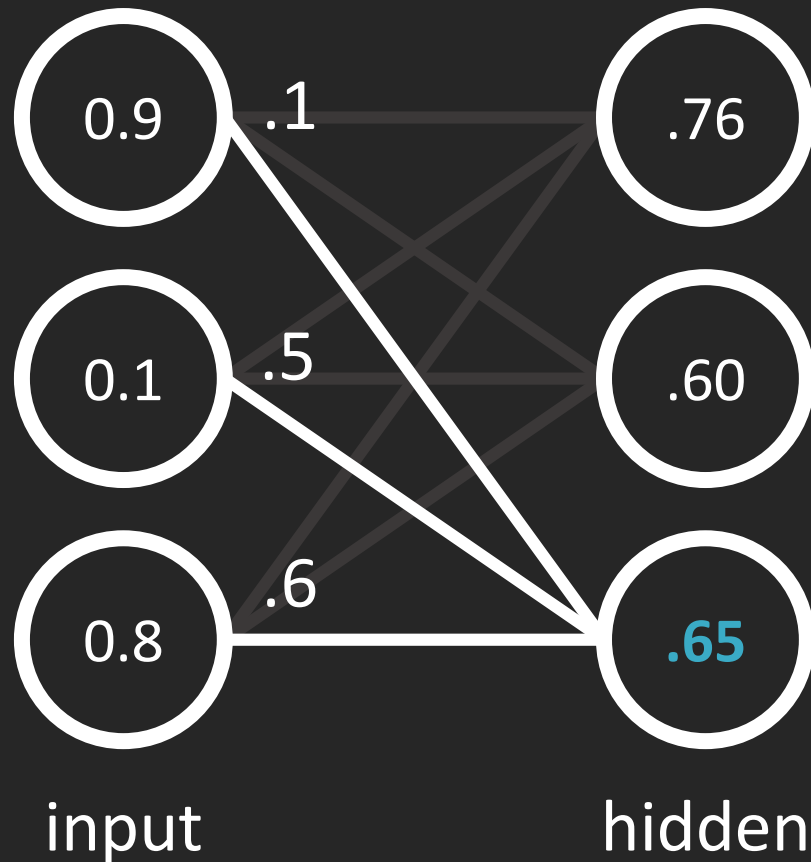
Weighted sum is passed into the activation function





ML Concepts


Activations become inputs for the next layer






ML Concepts

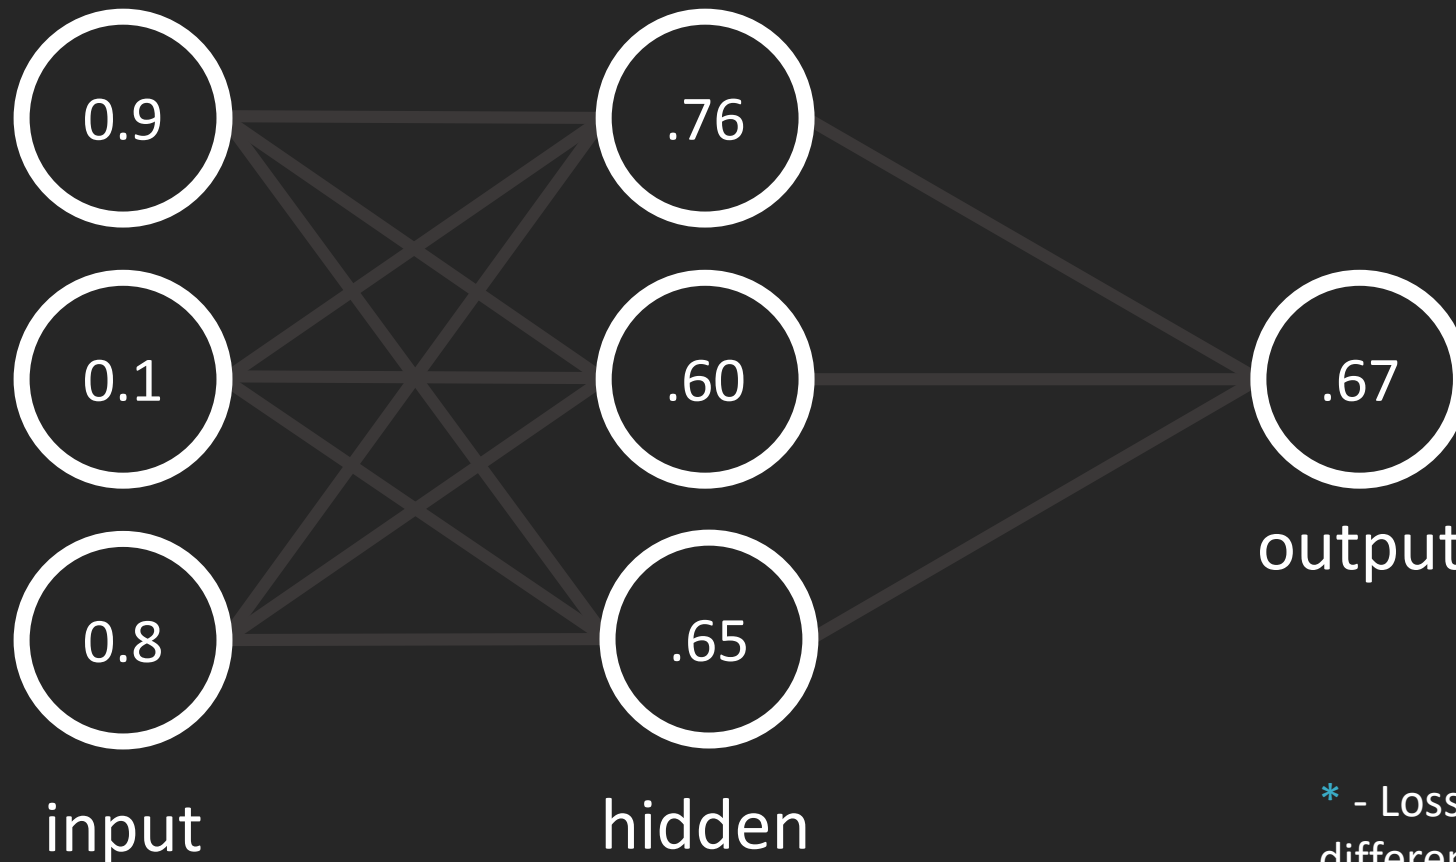
Output – Final result of cascading activations






ML Concepts

Loss - The gap between the target label and output *




target: 1
loss = 0.33

* - Loss can be calculated many different ways




ML Concepts

Backpropagation – Updating weights to improve * (“learn”)




* - Loss can be calculated many different ways





ML Concepts

Model – Network with learned weights



target: 1

~~loss = 0.33~~

loss = 0.25

Acc. + 0.08

Code – Neural Network

```
dataset = np.loadtxt(features.txt)
```

```
features = dataset[:,0:3]
```

```
labels = dataset[:,3]
```

```
model = models.Sequential()
```

```
model.add(layers.Dense(3, activation="relu", input_dim=3))
```

```
model.add(layers.Dense(3, activation="relu"))
```

```
model.add(layers.Dense(1, activation="relu"))
```

```
model.compile(loss="binary_crossentropy", optimizer="adam")
```

```
model.fit(features, labels, epochs=10, batch_size=10)
```



The Results

Total Samples: 130

Known Sandboxes: 18

Accuracy: 95%+

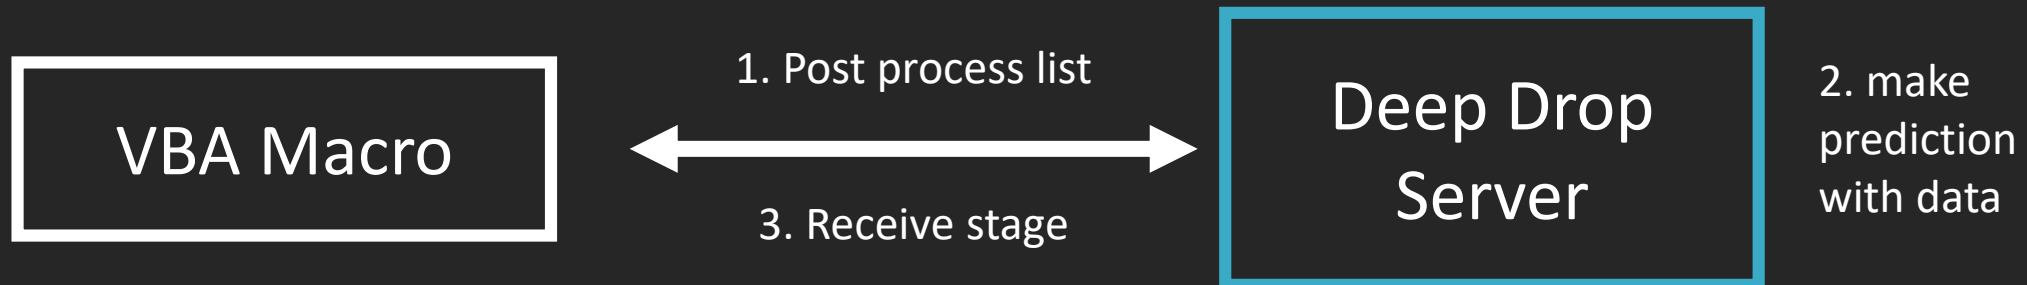
- Lots of options for feature combinations
- Ultimately a great problem for ML
- Doc2Vec / Decision Trees performed well too




Deep Drop - ML-enabled dropper server

<https://github.com/MooKitty/SchemingWithMachines/>

- Initially released at BSides LV 19
- Holds a trained model for parsing outputs
- Written in Python + Flask + Keras
- Makes stage delivery decisions based on ML



Data Pipelines





Data Collection

Issue: Offenders traditionally don't keep data

- Nobody is talking about this issue
- Should this remain an expectation?
- How can/should we anonymize data?
- Why are vendors allowed to keep data? (then sell it back)
- **Would sharing datasets further our field?**

Potential Solutions:

- Keep only model weights
- Hash features for storage
- Use models that don't require previous data



Textual Data

- Defenders already work with parsed data
 - SIEM collection and alerting
 - E-Mail report buttons
 - AMSI integrations
 - Known environments
- Offenders are still in a textual world
 - Command line interfaces
 - Screen/Session logs
 - ASCII art for every tool
 - Reports and narratives



Feature Engineering

- Target important meta-properties – **domain knowledge**
- Store as much as possible & build features later
- Use data analysis to assist
 - “What does the distribution of commands look like?”
 - Do new features line up with existing labels?
 - Reduce correlated features - “noise”
 - Once trained, find features which don’t affect outputs

Essentially ... Lots of **basic statistics**



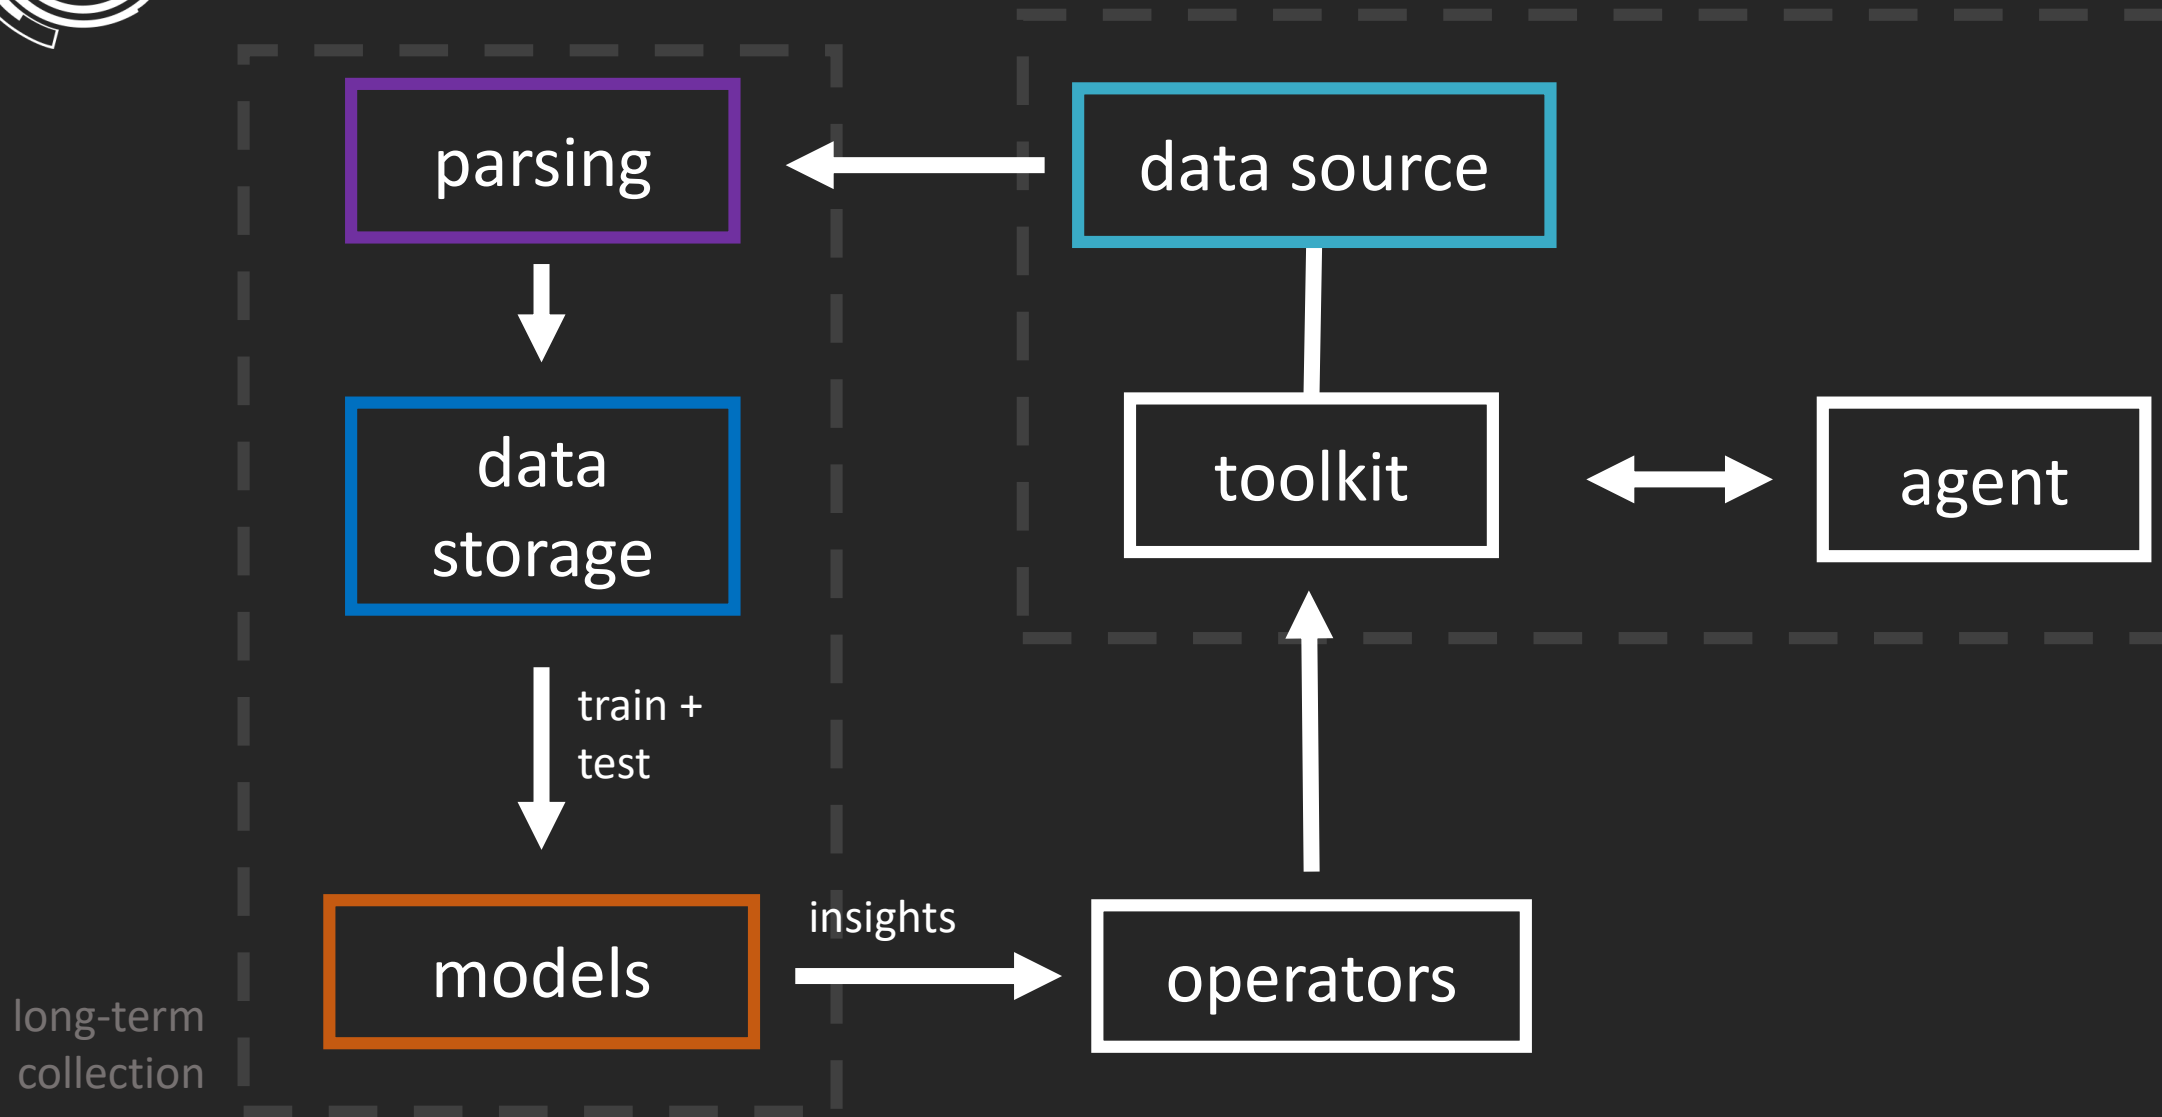
Data Pipelines

- Need systems for managing our data – long / short term
- This processing requires engineering
 - Start early to learn later
 - Implementations will vary by team TTPs
 - Solution will likely tie us down – agnosticism
 - Ideally passive – not interfering with ops
- Focus on high impact data
- Previous works on the subject:
 - https://github.com/ztgrace/red_team_telemetry
 - <https://github.com/outflanknl/RedELK>
 - <https://github.com/SecurityRiskAdvisors/RedTeamSIEM>



Data Pipelines

Op-period
processes





Team Integration

- Development is a requirement
 - Helpful if you also have/modify tool source
 - 2019 is the “year of C2” – shouldn’t be a problem
- Identify isolated jobs to begin delegating
 - Basic classification that a human already does
 - Suggestions that can augment decisions
- Basic statistics for ops
 - Average number of actions per operation
 - Count/distribution of commands and arguments
- Consider trust in the final solution

A composite image showing a view of Earth from space, with a bright, glowing horizon line. In the upper right, a large, reddish-orange sphere representing Mars is visible. The background is a dark, star-filled space. The text "Into the Unknown" is overlaid in the lower left.

Into the Unknown

State of Attack Paths

- Path finding with information is solved

Neo4j + Bloodhound

- **However, information will degrade**
 - Changes to Active Directory / Windows
 - Growing use of *nix in business
 - Network segmentation improvements
- **Networks are unknown, but discrete**
 - We **don't know** the user names, permissions, etc.
 - We **know** it's not infinite



Data Inference


Networks **require** data organization
therefore

Networks **imply** data organization



- Networks (AD) generally use text-labeling
 - We're all human, we expect it to be relational
- Can we infer these relationships?
 - Textual similarities
 - Mapped drives, local users, host information
 - LinkedIn, GitHub, public exposure ...




A Network of Probabilities



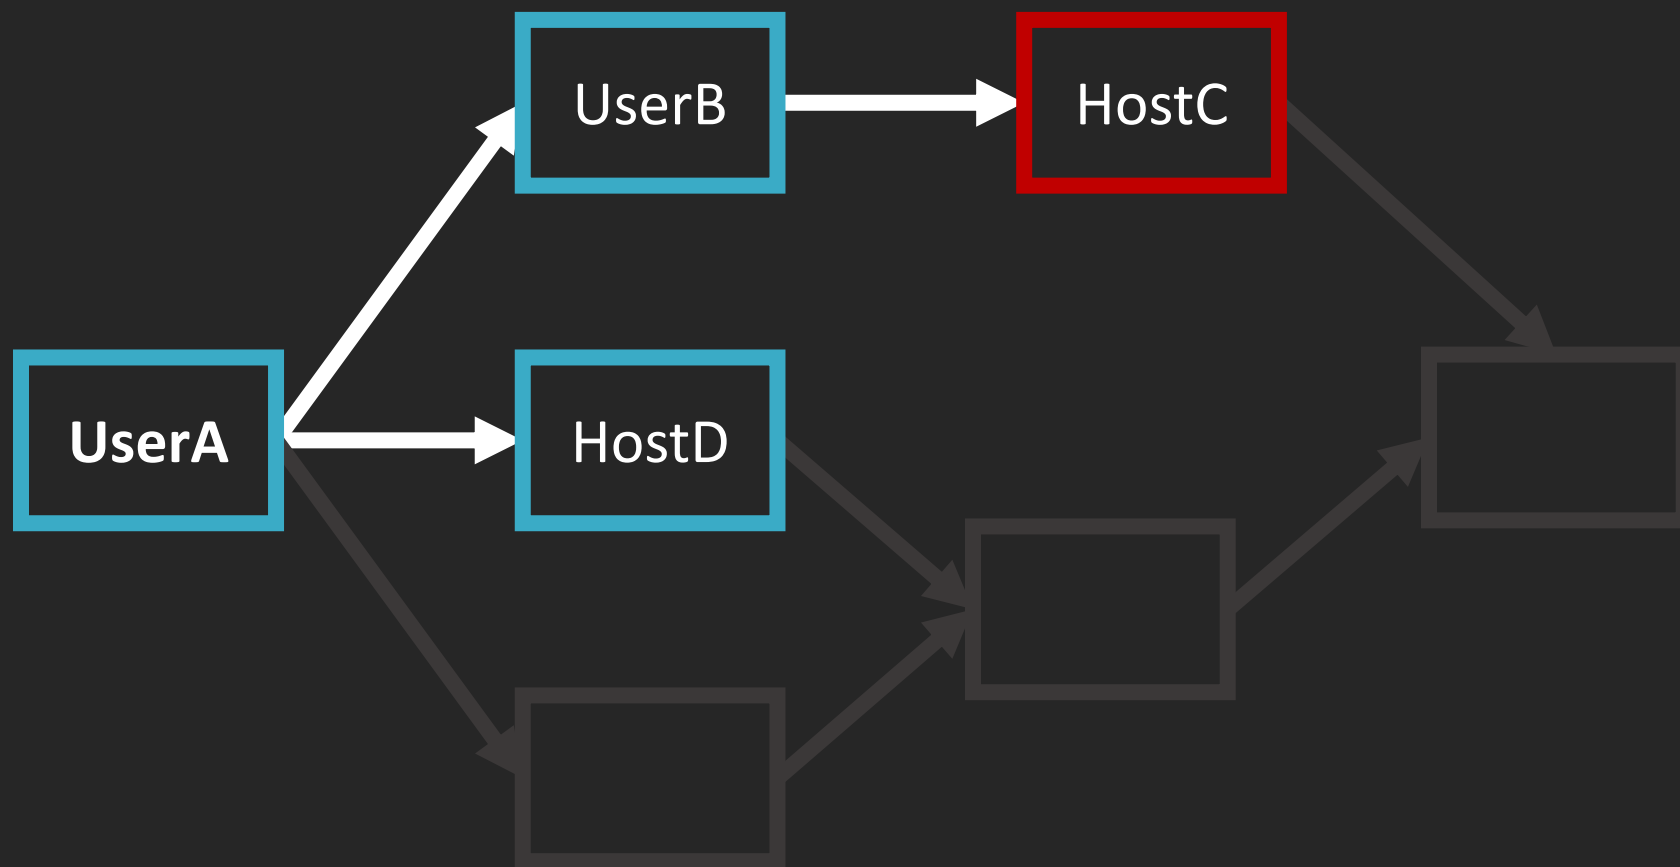
A Network of Probabilities



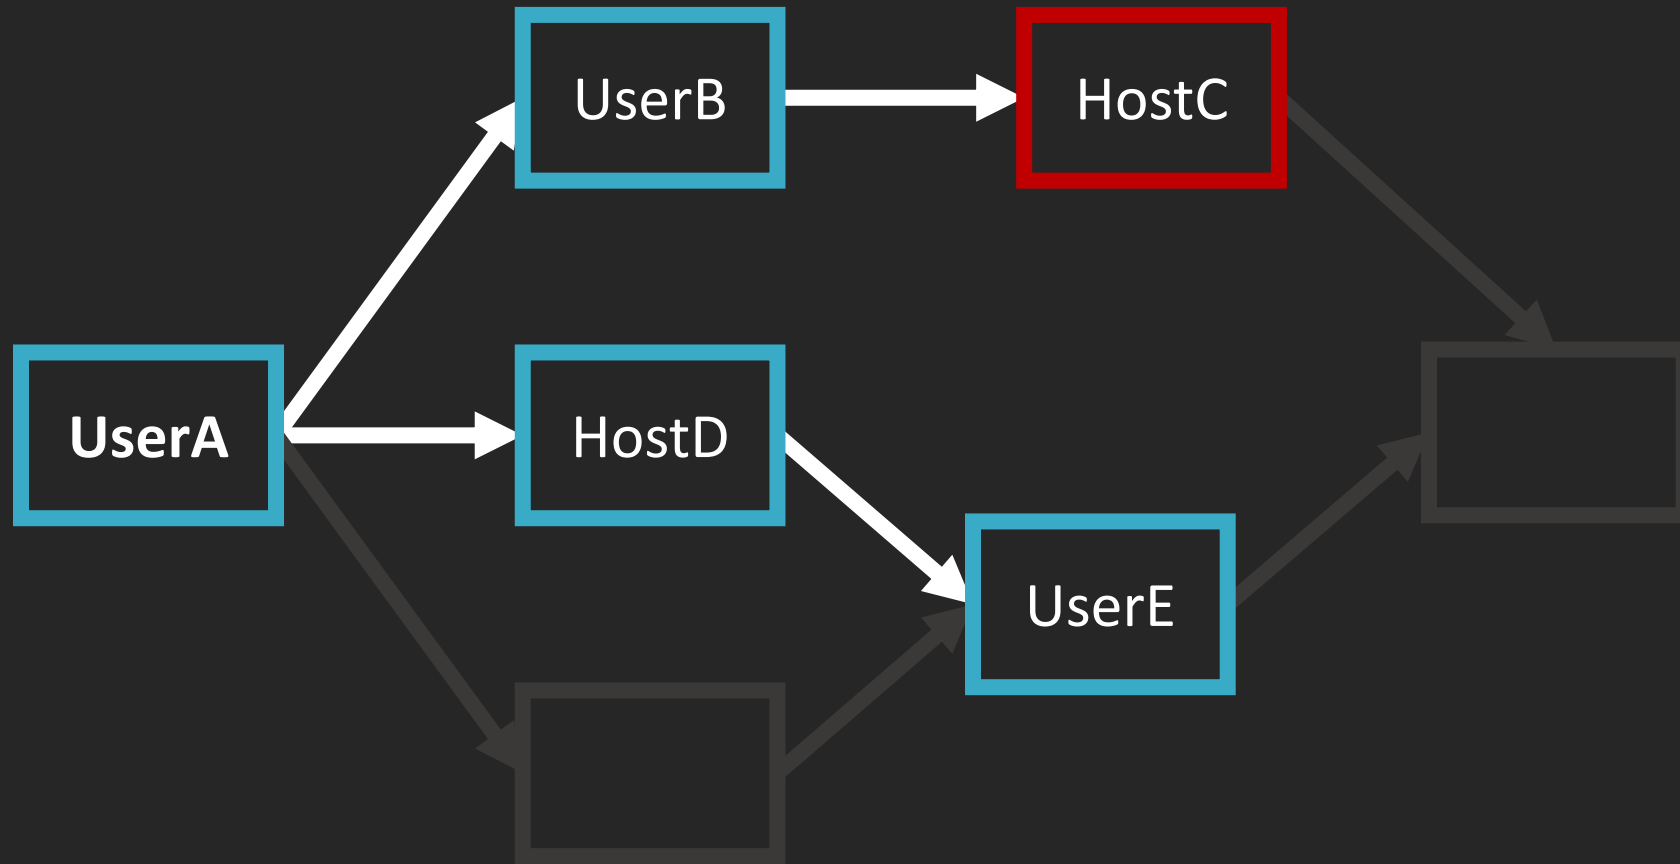
A Network of Probabilities



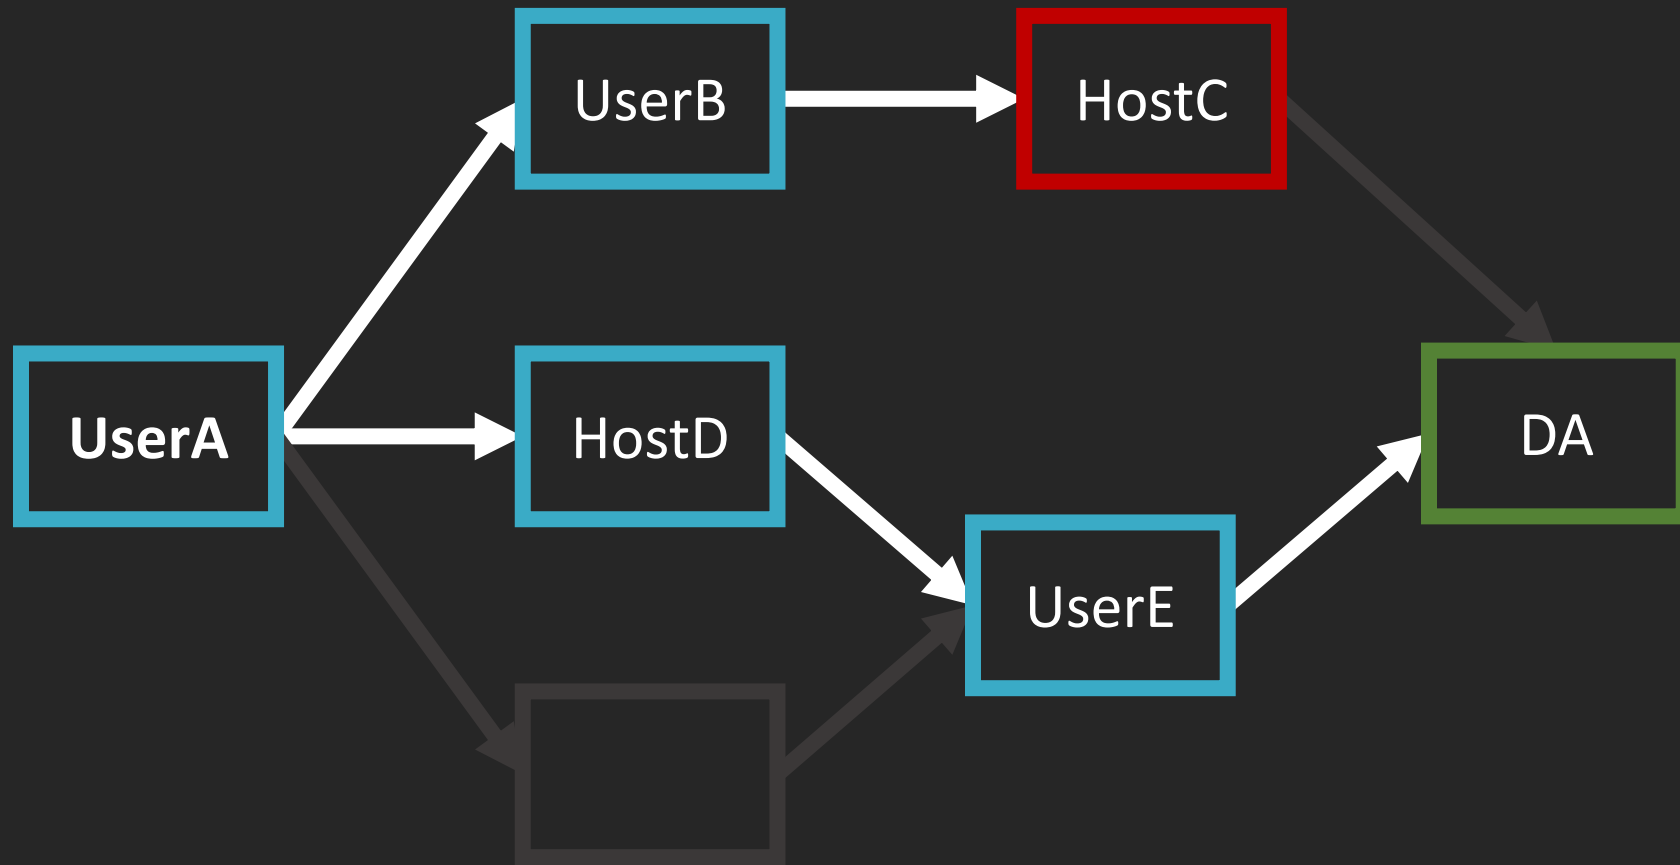
A Network of Probabilities



A Network of Probabilities



A Network of Probabilities



Mental Models

As operators, we build mental network maps

- We assume relationships based on:
 - Standardized textual labels
 - Experiences in the network itself
 - Pattern recognition – how **has** it been configured, how **will** it
- We act on these assumptions with queries
 - Validate access to a host
 - Verify the group membership of a user
 - Collect new attack surface via enumeration

Information creates confidence



Simulating Mental Models

Data + Heuristic Search + Simulation

- **Data** - Information from the current context
- **Heuristic** - Relationships between data points
 - Operator driven - flexible
 - We can “assume” relationships, or even new data points
 - Use algorithms to bring up the most relevant data
- **Simulation** - Select actions based on heuristics
 - Could include the “impact” of potential actions
 - Can assist an operator, or **become one** – active / passive



Data Layer

- Host/User/Group name information
- Active network connections
- Outbound RDP history
- Network queries
 - Active Directory with limited filters
 - Direct host service access
- Host-based Events
 - Event logs
 - ETW tracing
 - Custom tracking over time



Heuristics Layer

- How can we relate textual data?
- What strategies are used already?
- What would an operator care about?

We need some number to support our simulation



Heuristic: Simple 'if' Statement

- Operator driven insights
- Doesn't require complexity

```
if <output> in command_output:  
    return new_state, reward  
else:  
    continue
```




Heuristic: Cosine Similarity

```
match (g:Group)
with collect(g) as groups
match (u:User)
with u, algo.ml.oneHotEncoding(groups, [(u)-[:MemberOf]->(memberof) |
memberof]) as embedding
with {item:id(u), weights: embedding} as userData
with collect(userData) as data
call algo.similarity.cosine(data, {similarityCutoff: 0.7071, write: true, topK:
100})
yield nodes, p50, p75, p90, p99, p999, p100
match (u:User {name: "<USER>"})-[similar:SIMILAR]->(other)
return other.name as user, similar.score as score
```



Heuristic: Cosine Similarity



Heuristic: Levenshtein Distance

```
match (c:Computer)
match (g:Group)
with g,c, apoc.text.levenshteinSimilarity(g.name, c.name) as data
return g.name as group, c.name as host, data as score
order by score desc
limit 100
```



Heuristic: Levenshtein Distance

Group	Hostname	Score
SQL Developers	SQLDEV01	55
Domain Admins	DOM-PRINS	55
SQL Admins	SQLDEV01	44
VPN Users	DEVVPN-B	38



Simulation: Shortest Path

Dijkstra's Algorithm

- Shortest distance from A->B
- See their awesome talks - the past 5 years
- Useful for map/path finding problems
- Currently what Bloodhound uses
 - not the only one we could use ...

<https://github.com/andyrobbins/PowerPath>



Simulation: Shortest Path

A-Star Algorithm

- Shortest distance from A->B + **Heuristic**
 - Helps us avoid particular paths
 - Ignore paths which might be unavailable – segmentation
 - Punish “noisy” paths

```
CALL algo.shortestPath.astar.stream((startNode:Node, endNode:Node,  
weightProperty:String, propertyKeyLat:String, propertyKeyLon:String,  
  {nodeQuery:'labelName', relationshipQuery:'relationshipName',  
direction:'BOTH', defaultValue:1.0})
```

```
YIELD nodeId, cost
```



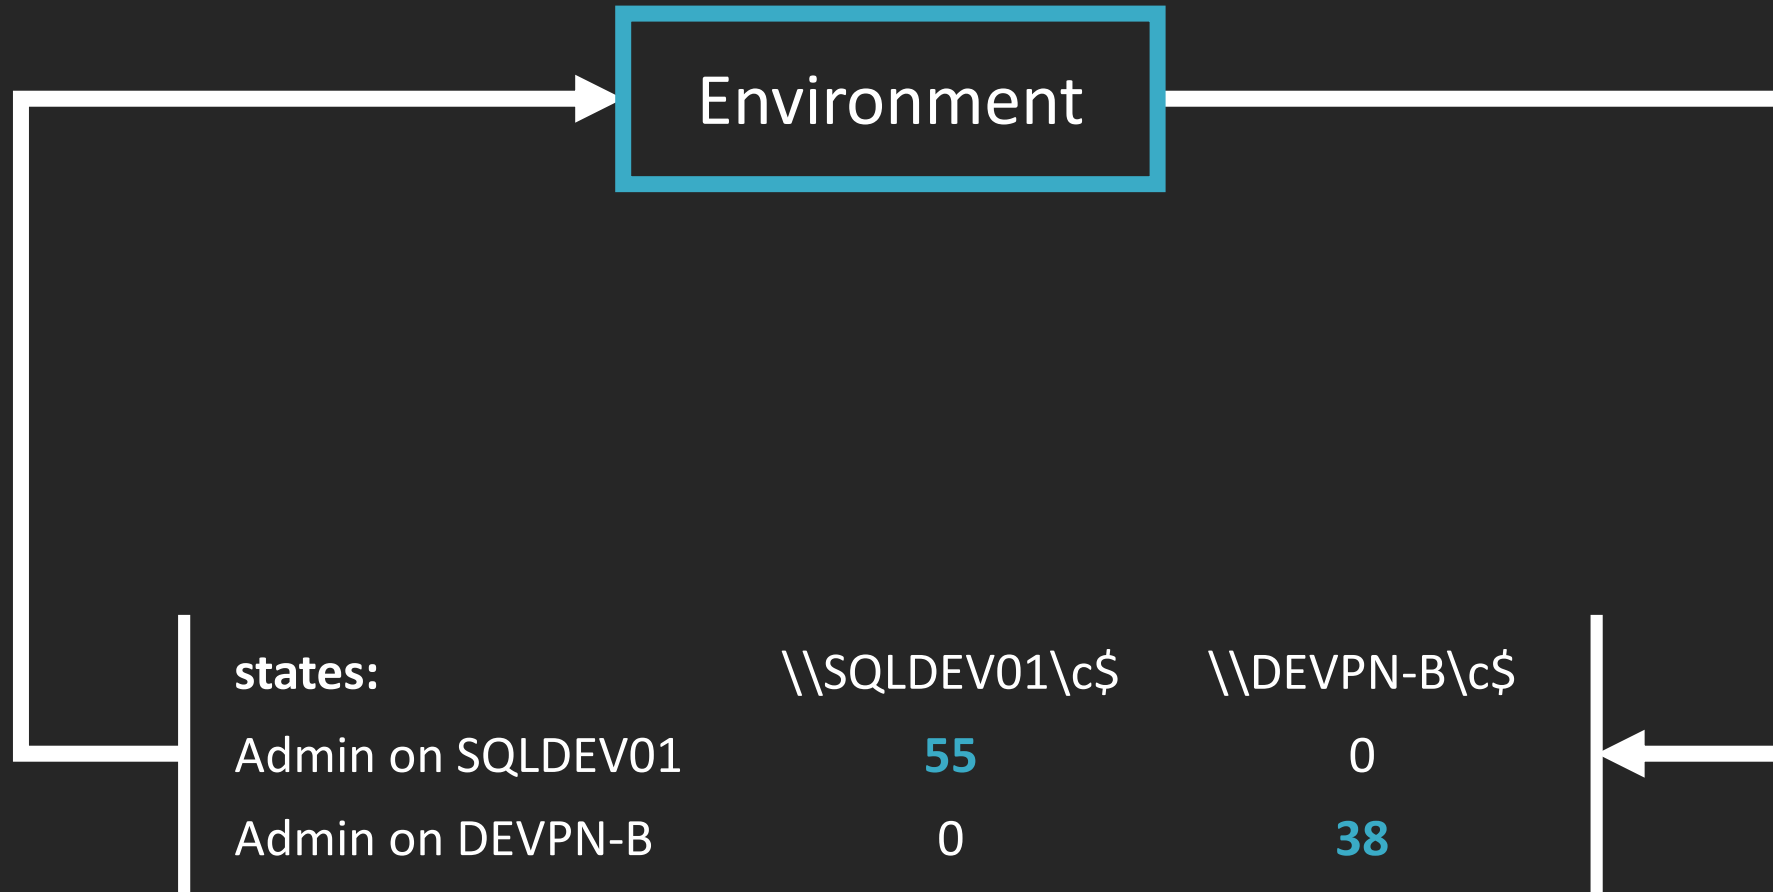
Q-Learning for Automation

For a given action in a given state, the environment returns a new state, and a reward

- Basic reinforcement learning
- Allows an agent to learn optimal actions
- Strength is through the heuristic you use
- Requires some initial state – all weights are 0

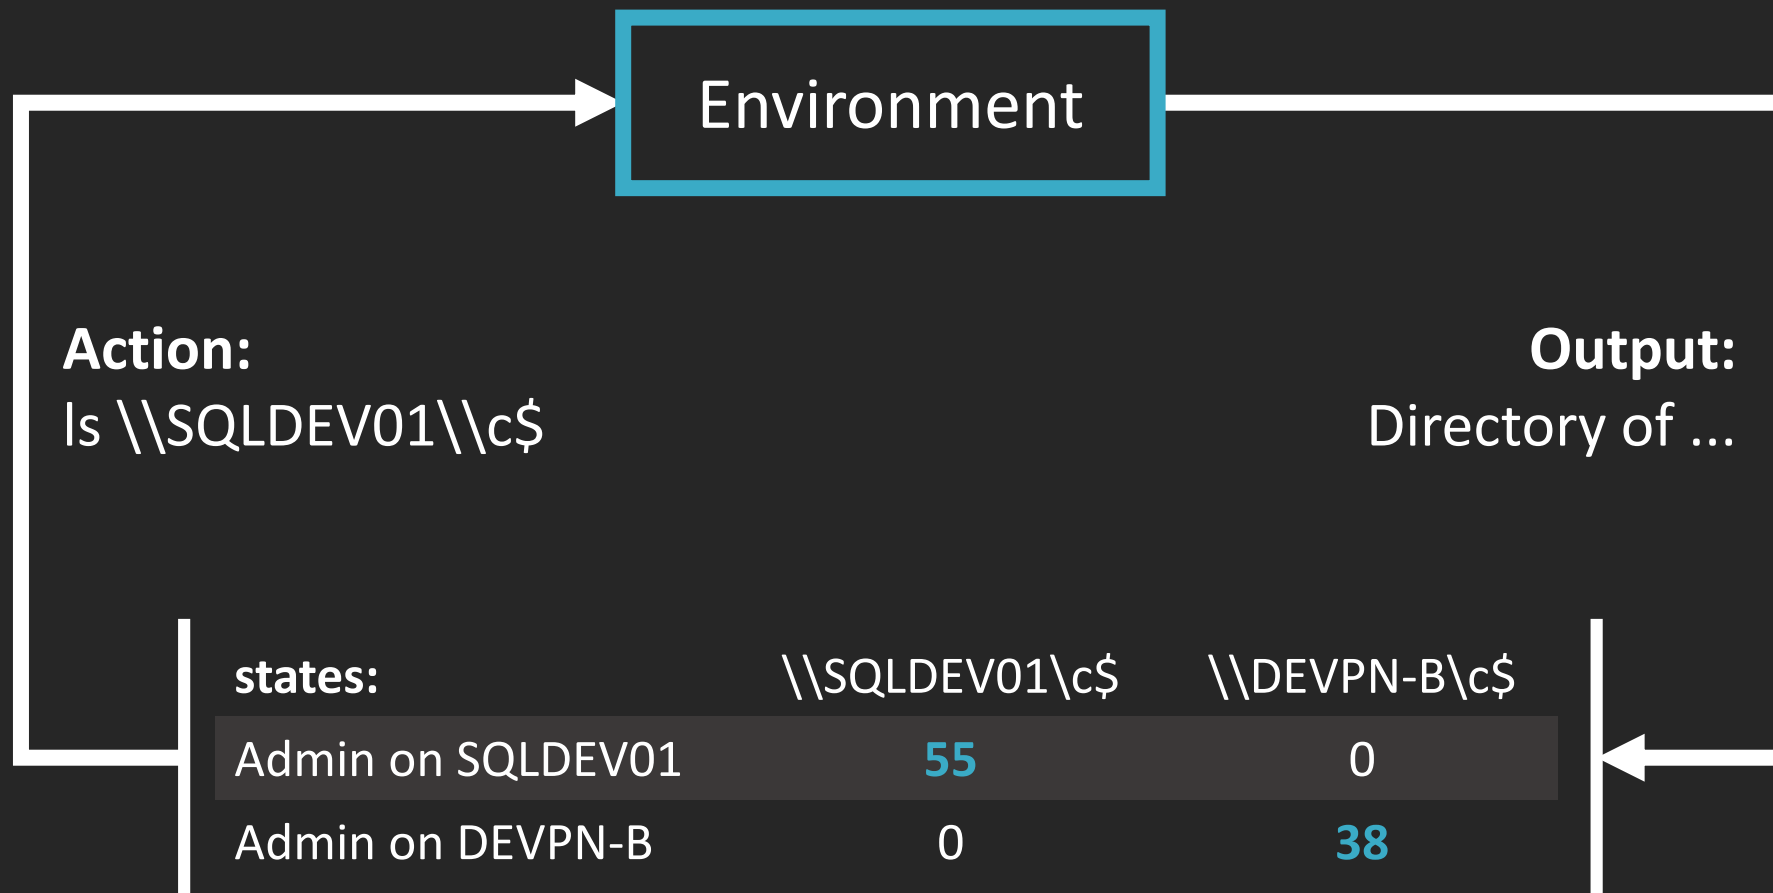


New State: ?
Reward: ?



New State: Admin on SQLDEV01

Reward: +55



New State: Shell on SQLDEV01
Reward: +500



Back on the radar




Goal > Data > Prior > Query > Posterior

- We are all just data processors
- Attack graph theory is the future
 - Bloodhound/resiliency is like white-box code review
 - **Heuristic-based cyclic queries are black-box**
- With limited knowledge, we work in probabilities
- Given a sufficient process, Q-learning could op



Adversarial Attacks



Adversarial ML

“Attacking existing models”

Effective | Efficient vs. Secure | Robust

- Proven mainly in the lab
 - Not theoretical, just hard to build
 - Many demonstrations lack practical use
- Two basic approaches:
 - **White:** Access to the original model, architecture, etc.
 - **Black:** Access only to the outputs for a given input



Previous Works

- **DeepWordBug** - Black-box generation of adversarial text

<https://github.com/QData/deepWordBug>

- **Cylance, I kill you!** - Client-side model reversing

<https://skylightcyber.com/2019/07/18/cylance-i-kill-you/>

- **Good word attacks on statistical mail filters**

<https://ix.cs.uoregon.edu/~lowd/ceas05lowd.pdf>

- **Robustness Toolbox** – Attacks, defenses, etc.

<https://github.com/IBM/adversarial-robustness-toolbox>



proofpoint® - Case Study

- **E-Mail security company**
 - Spam detection
 - Malware sandboxing
 - URL analysis
 - End user training
- Openly promote their use of ML (MLX, CLX)
- Supporting 230k+ domains – Rapid7 Sonar DNS data
 - 590 **gov** domains
 - 2300 **edu** domains



proofpoint® - Vulnerability

To: <reciever@domain.com>
From: <sender@domain.com>
Subject: Our Meeting

...

X-Proofpoint-Spam-Details: rule=nodigest_notspam policy=nodigest score=0
malwarescore=0 mlxlogscore=999 mlxscore=0 suspectscore=14 spamscore=0
impostorscore=0 adultscore=0 clxscore=593 priorityscore=0 phishscore=0
bulkscore=97 lowpriorityscore=97 classifier=spam adjust=0 reason=mlx
scancount=1 engine=9.1.0-12345000 definitions=main-12345

Leaky inputs... tsk tsk



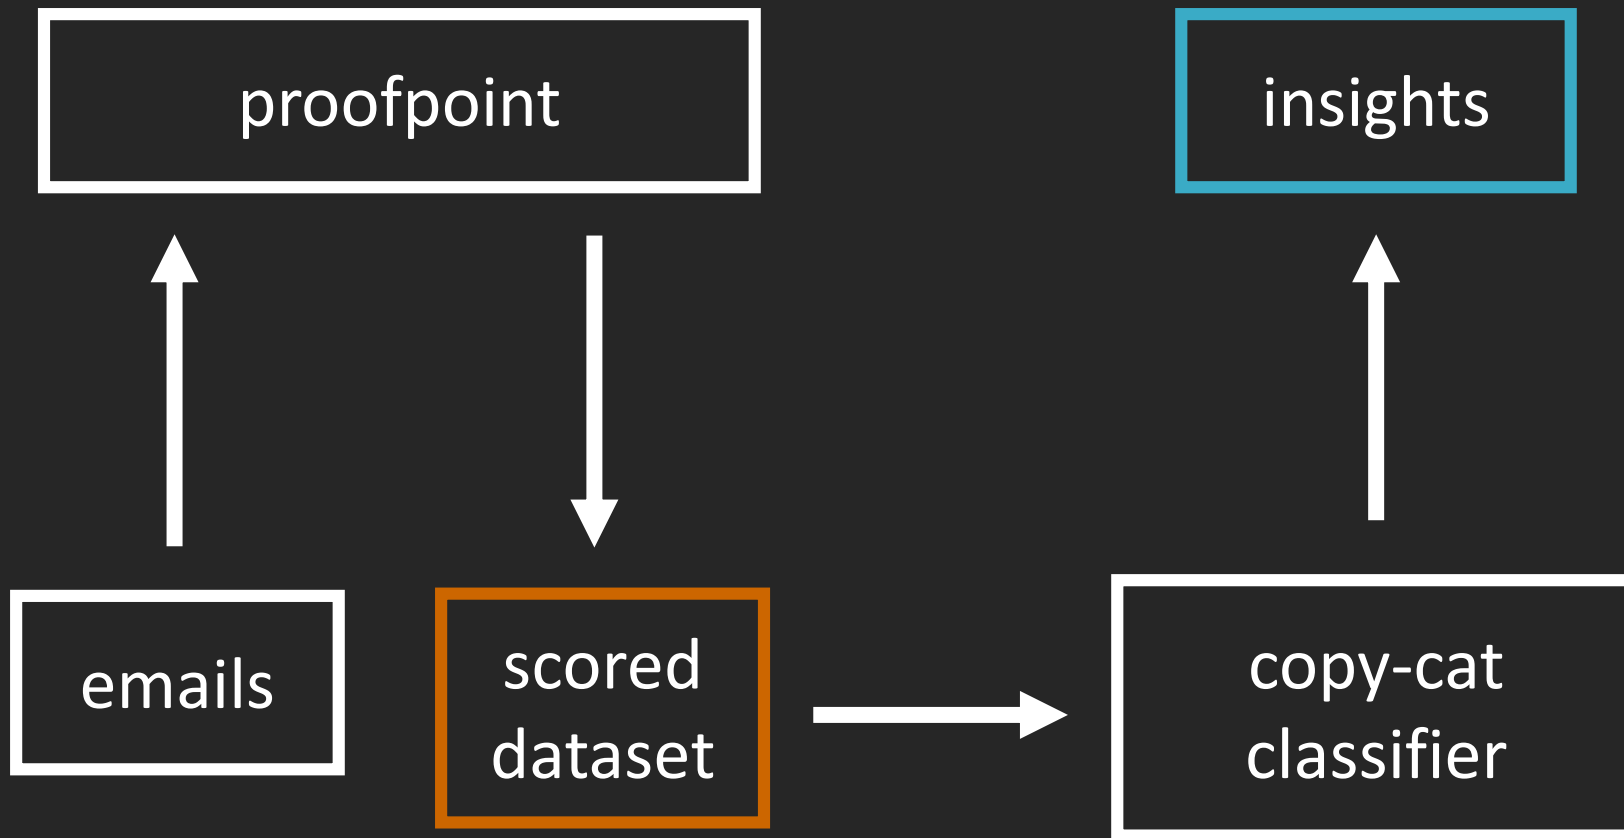
proofpoint® - Attack (a)

1. **Collect a dataset** – Send X emails to steal scores
2. **Copy the model** - Use their outputs to duplicate
- 3a. **Extract information from the model**
 - Take N highest/lowest emails, unique the words
 - **Toggle inputs to discover the most impactful tokens**
 - Invert the model mathematically *
 - Randomly alter/add content and re-score
(char swaps, homoglyphs, tense)

* - <https://r2rt.com/inverting-a-neural-net.html>



proofpoint® - Attack (a)

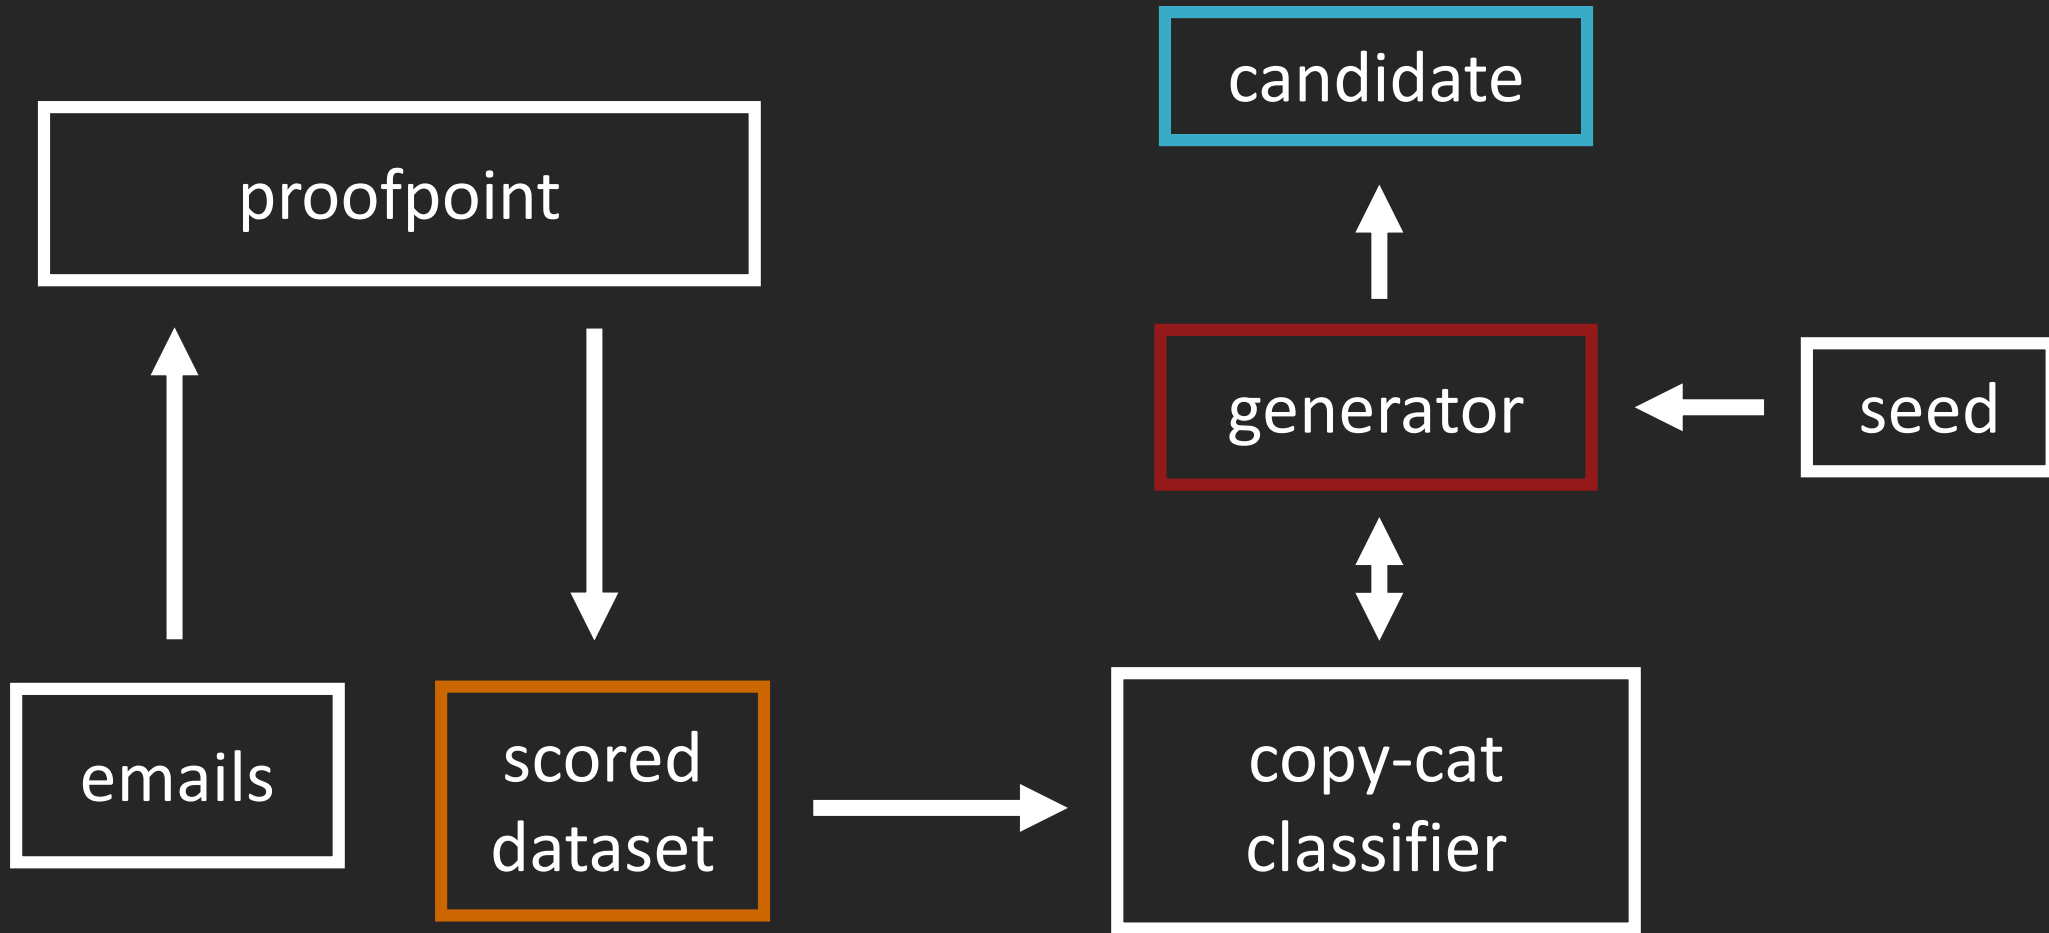


proofpoint® - Attack (b)

1. **Collect a dataset** – Send X emails to steal scores
2. **Copy the model** - Use their outputs to duplicate
- 3b. **Make a generator** – Use our copy-cat to train
 - Let it learn useful insights
 - Target maximum score – custom loss function
4. **Automate improvements**
 - “Fix” pre-written candidates
 - Generate “good” content from scratch



proofpoint® - Attack (b)



proofpoint® - Challenges

- **Initial email content**
 - Finding a sufficient dataset
 - Links/attachments have large effects on the score
- **Bulk email delivery**
 - Easy if we have a Proofpoint inbox – harder if we don't
 - Total emails required for a sufficient dataset
 - Extraction process altering the scores
- **Final Outcomes**
 - Generators will likely create gibberish – human intervention
 - Bypassing a model is only one part of the “defenses”



1. Collect a dataset

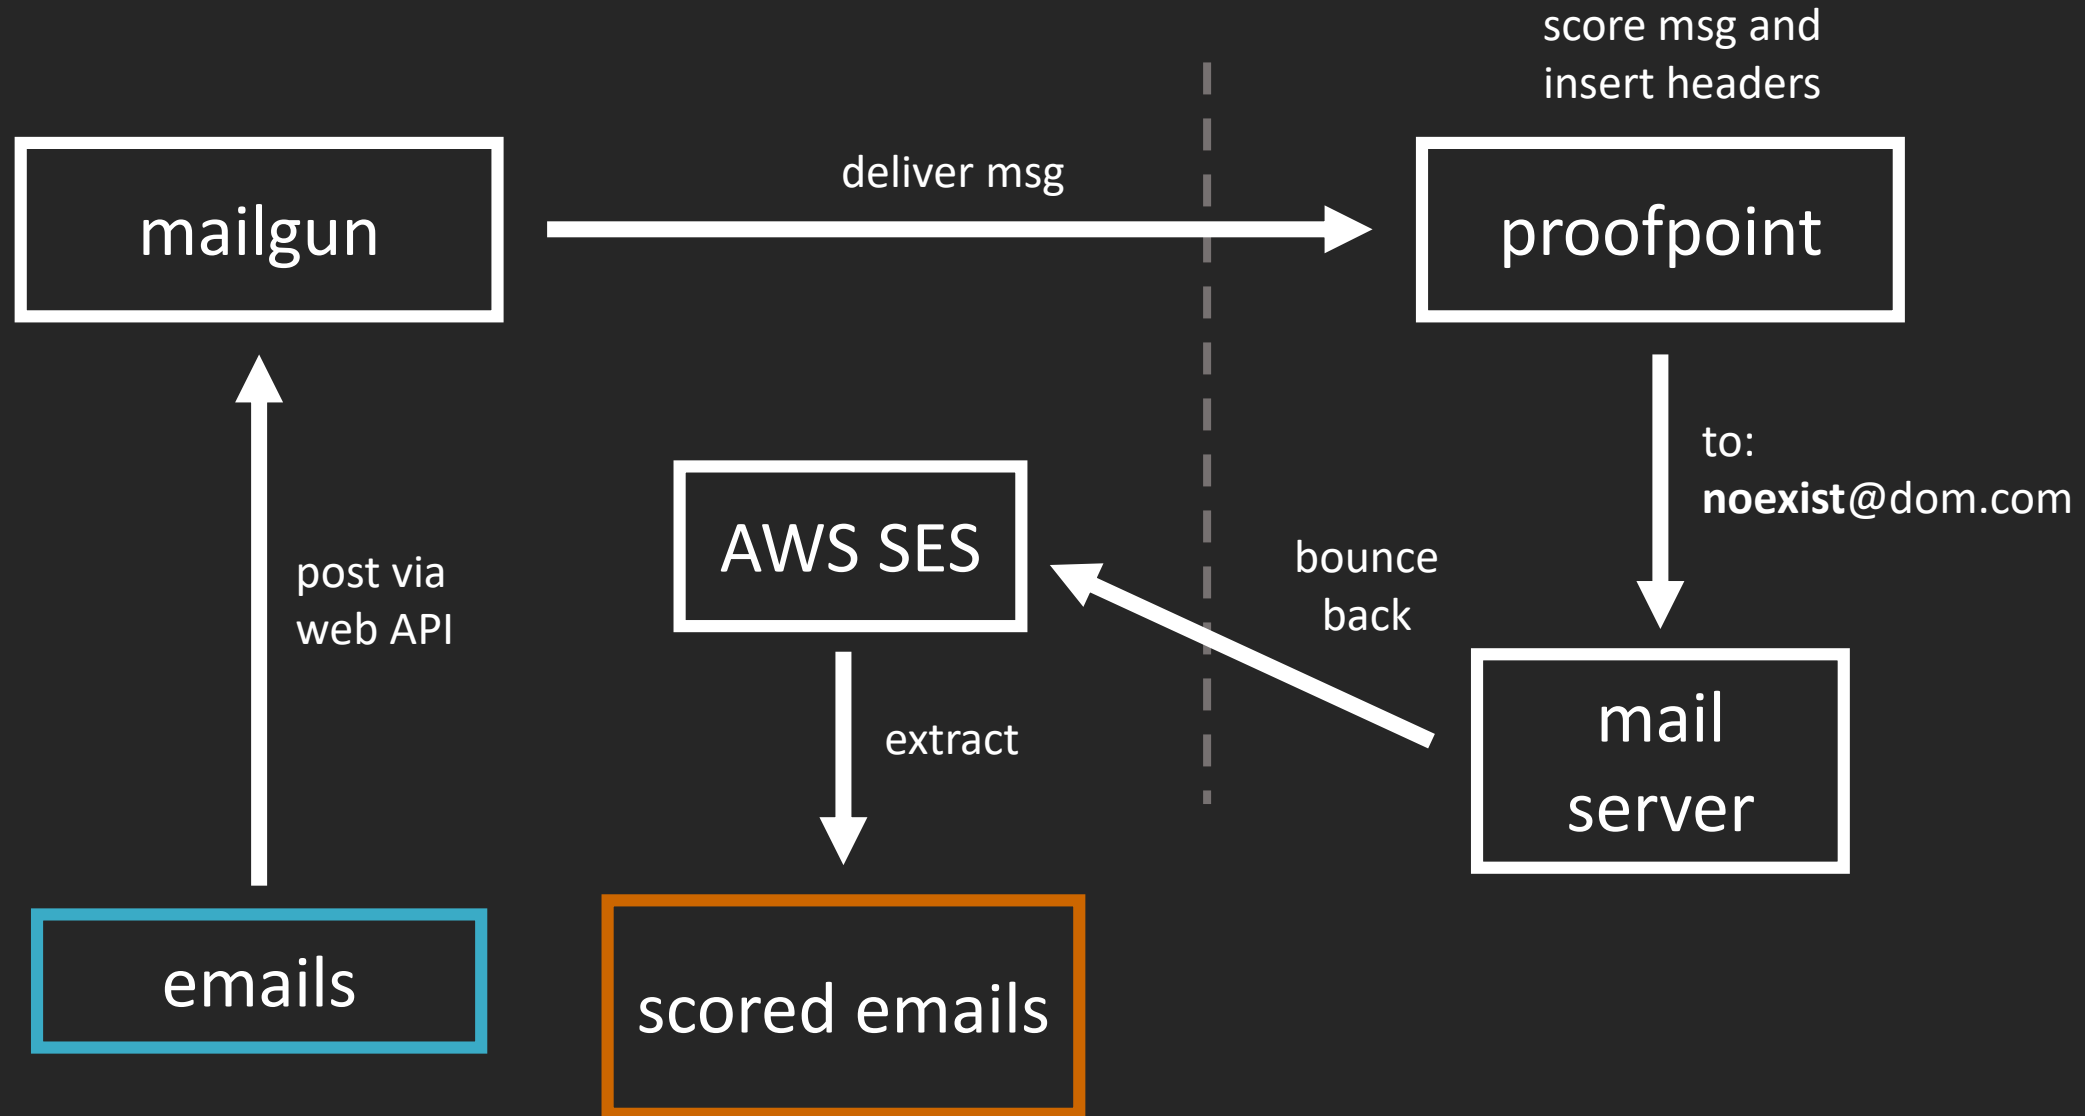
- Needed to gather inputs for scoring (a lot)
 - Enron dataset for text-based candidates
 - ISCX-URL-2016 for link-based candidates
- Use bounce-backs to collect the scores
 - Delivered using Mailgun
 - Received using AWS SES + S3 bucket

We ran multiple collection runs:


1. 5k pre-processed/scored samples from Enron
2. 13k Links inside a generic template
3. 15k raw subject + bodies from Enron inboxes



1. Collect a dataset



1b. Data Analysis



score	bulk	malware	priority	spam	phish	impostor	mlx	mlxlog	low-p	suspect	adult	clx
16	0	0	90	16	0	0	16	73	0	3	0	403
0	0	0	118	0	0	0	0	505	0	8	0	324
0	0	0	118	0	0	0	0	479	0	19	0	303
0	0	0	118	0	0	0	0	489	0	3	0	315
0	0	0	118	0	0	0	0	538	0	3	0	321
0	0	0	90	0	0	0	0	605	0	3	0	437
0	0	0	118	0	0	0	0	455	0	3	0	293
0	0	0	90	0	0	0	0	728	0	3	0	466
0	0	0	118	0	0	0	0	477	0	3	0	299
0	0	0	118	0	0	0	0	483	0	3	0	288
0	0	0	118	0	0	0	0	484	0	3	0	344
0	0	0	90	0	0	0	0	595	0	74	0	432
0	0	0	118	0	0	0	0	329	0	3	0	304


1b. Data Analysis

	<i>score</i>	<i>bulk</i>	<i>priority</i>	<i>spam</i>	<i>phish</i>	<i>mlx</i>	<i>mlxlog</i>	<i>low-p</i>	<i>suspect</i>	<i>adult</i>
bulk	-									
priority	-	-								
spam	1	-	-							
phish	-	-	-	-						
mlx	1	-	-	1	-					
mlxlog	-.2	-	-	-.2	-.1	-.1				
low-p	-	1	-	-	-	-	-			
suspect	-	-	-	-	-	-	-	-		
adult	-	-	-	-	-	-	-	-	-	
clx	-	-	1	-	-	-	-	-	-	-

* - values < .1 have been omitted


1b. Data Analysis

15k **text**-based samples
mlxlogscore



1b. Data Analysis


13k **link**-targeted samples
mlxlogscore



1b. Data Analysis

10k **link**-targeted samples
mlxlogscore

Activation
function
baby!



2. Copy the Model

- Select a label for training: **mlxlogscore**
 - Good distribution – at least for links
 - Previously scaled / activated
 - *Generally* between 1 and 999
 - Larger = "safer"
- Select some likely model emulators
 - **Neural Network** + Bag of Words (BOW) *
 - **LSTM** + Sequenced Text

* - Also referred to as "one-hot" encoding



2. Copy the Model - Results

	Neural Network + BOW	LSTM + Sequences
Text targeted samples	69	91
Link targeted samples	42	96

(we didn't plan this, we swear)

* showing scaled mean absolute error (mean **point** error)



3. Extract Information

- **Make text alterations and rescore**

- Take a phishing email: “Click here for cats”

- **Clikc** here for cats (Swap)

- Click hare **fur** cats (Substitute)

- Click **her** for cats (Delete)

- Click here for **carts** (Insert)

- Final outputs need to make sense

- **Toggle input tokens and rescore**

- [1,1,1,1] - Click here for cats - 500

- [0,1,1,1] - here for cats - 490

- [1,0,1,1] - Click for cats - **300**

- **Score every possible combination of tokens - fuzzing**



3. Extract Information

```
for sample in test_set:
    base = make_prediction(sample)

    for token in sample:
        altered = sample.toggle(token)
        test = make_prediction(altered)

        # Record a rolling score movement
        insights[word] += (base - test)
```



3. Extract Information - Texts

good

calculation

asset

appreciated

finalized

tyson

difficult

dial

default

lawyers

bids

meh

lisa

digest

piano

stems

architectual

living

smells

storms

alcoholic

broccoli

bad

software

99

unsub

bridgeline

absolutely

quantities

hydro

proposal

deposit

holden



3. Extract Information - Links

good

movies
category
ecnavi
xml
payment
docs
shop
dest
kitchen
webapps

meh

cpanel1
certificate
area2
delores
verify2
struggles
chinas
second
webserver
uniq

bad

cool
citi
hc
wp
license
includes
styles
logon
plugins
spreadsheet



Confirming our insights - **Texts**

Top 10 **highest** scoring words:

999

Random 10 words from the middle:

640

Top 10 **lowest** scoring words:

...

**mx0a-000a1001.pphosted.com gave this error:
This message looks too much like SPAM to accept.**



Confirming our insights - Links

<https://neverexistdomain.com/wp-includes/file>

Predict: 378 | Real: **300**

<https://neverexistdomain.com/up-uncludes/file>

Predict: 600 | Real: **559**

<https://.../ecnavi/category.xml?movies=payment>

Predict: 999 | Real: **999**



Disclosure | Remediation

- **Models are interesting beasts**
 - Represent learned vectors – not always apparent
 - Difficult to retrain / rebuild
 - Black box with “magic” inside
- **What warrants responsible disclosure?**
 - What % is considered a viable bypass?
- **How does remediation occur?**
 - Can't just add a signature
 - General models might work even without leaky outputs



Final Thoughts



Real World Talk

- **Application Whitelisting**
 - Was cool until people realized there were bypasses
 - Will be a vendor pitch while it gets sales
- **EndGame ML Competition**
 - Simple bypasses for static analysis (sRDI, emojis)
 - Data scientists solving defensive problems
 - Static analysis is only a small part of the battle


“If you don’t understand X before ML, you won’t understand it after” - Will

Not All Bad ...

- **The next generation of malware**
 - Intelligent agents
 - Genetic programming client-side (JIT, variants on the fly)
 - Distributed API calls and hooks
 - Rootkits – layered defenses (clothing)
- **Securing and hardening models**
 - Helping vendors improve their products
 - Ensuring ML isn't the next big security gap



Fun Projects



- What other defenses leak outputs?
 - Windows defender AMSI sampling
 - URL categorization
- Can we evade other mechanisms?
 - E-Mail attachments in transit
 - HTML content during site inspection
- What other offensive tasks can be offloaded?
 - File & directory enumeration
 - External reconnaissance
 - Chat/E-mail bot for phishing
- Add data extraction to Seatbelt for Neo4j

Final Thoughts

- Lots of fun work to be done, come play!
- Machine learning is here to stay, don't sleep on it
 - It's all a joke until you get caught
 - "Model bypasses" will become a part of offensive kits
 - ML understanding could quickly become a job requirement

**"Might be nothing – could be everything.
Likely somewhere in the middle.
Time will tell." - Will**



Greetz

@culteredphish – Colleague & all around good guy

@tyler_robinson – Ex-Colleague & long time friend

@rharang – Answered some helpful questions

@silentbreaksec – Company supporting this research

- Nancy Fulda of BYU
- Will's Mother-in-law for babysitting

All of you for attending the talk



Find Us After



Will Pearce
@moo_hax
MooKitty



Nick Landers
@monoxgas



SILENTBREAK
SECURITY

Soon: <http://github.com/MooKitty/FourtyTwo>

Resources

- **“Make your own Neural Network”** – Tariq Rashid
- **“3 Blue 1 Brown”** YouTube channel
- **“Jabrils”** YouTube channel

<https://www.kaggle.com/>

<https://silentbreaksecurity.com/machine-learning-for-red-teams-part-1/>

<https://github.com/MoooKitty/RedML>



So long and thanks for all
the phish!

